

NIHAR GUDADHE

Data Scientist | AI/ML Engineer

Nagpur, India | +91- 9168844536 | gudadhenihar@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Scientist and Analyst with 2 years 8 months of experience across the full machine learning lifecycle from EDA and feature engineering to model development and production deployment, for regression, classification, and time-series problems. Built and deployed a Random Forest pricing model behind a FastAPI service running in Docker that cut quotation turnaround from days to seconds. Automated CRM and ERP data pipelines using MySQL, dbt, and PostgreSQL to streamline the Quote-to-Cash process, and delivered Power BI dashboards that improved throughput, conversion, and return rates. Currently expanding into Generative AI, including LLMs, RAG, LangChain, and prompt engineering, alongside MLOps practices such as MLflow and drift monitoring.

TECHNICAL SKILLS

Programming & Databases: Python, SQL, PostgreSQL, MySQL, Excel

Machine Learning & Statistics: Regression, Classification, Clustering, Time Series Forecasting, Ensemble Methods (Random Forest, XGBoost, LightGBM), PCA, Anomaly Detection, Feature Engineering & Selection, Cross-Validation, Hyperparameter Tuning, SMOTE, Hypothesis Testing, A/B Testing

Deep Learning & NLP: ANN, CNN, RNN, LSTM, GRU, YOLO, OpenCV, spaCy, Encoder-Decoder, Attention Mechanism, Tokenization, Word Embeddings, Named Entity Recognition, Sentiment Analysis

Generative AI: Large Language Models (LLMs), Prompt Engineering, RAG, Hugging Face Transformers, LangChain

Libraries & Frameworks: Pandas, NumPy, Scikit-learn, TensorFlow, PyTorch, Matplotlib, Seaborn, NLTK

Data Visualization: Power BI, DAX, Tableau

Engineering & Deployment: FastAPI, Docker, dbt, MLflow, Git, GitHub, CI/CD, Render, Jupyter Notebook, VS Code

Core Concepts: Statistical Inference, Confidence Intervals, EDA, Data Cleaning, Data Modeling, Data Warehousing, Data Wrangling, ETL/ELT, Model Evaluation & Validation, Overfitting & Regularization, Data Quality Assurance, Business Intelligence, KPI Design & Tracking, Data Storytelling, Stakeholder Reporting.

PROFESSIONAL EXPERIENCE

Data Scientist & Analyst | Influx Control Panel, Nagpur

Feb 2024 - Present

- Built an automated ELT pipeline that pulls raw CRM leads and billing records into a MySQL staging layer, transforms the data with dbt, and loads clean tables into a PostgreSQL warehouse, replacing the manual data pulls the team relied on before.
- Trained and deployed a Random Forest regression model in Scikit-learn to estimate costs for complex industrial items automatically. Running behind a FastAPI endpoint in Docker, it cut quotation turnaround from days down to seconds.
- Built a production line efficiency dashboard to track cycle times and machine idle time. The bottlenecks it surfaced on the shop floor led to a **12%** increase in throughput.
- Designed a 360-degree customer dashboard in Power BI covering account value, engagement, and conversion stage, which the sales team used to prioritize high-value accounts and improve conversion rates by **10%**.
- Applied anomaly detection to flag defects by correlating raw material batches with support ticket trends, helping engineering catch recurring issues earlier and cut product return rates by **14%**.

RELEVANT EXPERIENCE

Business Operations Lead | Hardware Business, Amravati

Mar 2020 - Nov 2023

- Managed end-to-end inventory operations for a family-owned hardware business, including stock planning, procurement scheduling, and vendor coordination.
- Tracked sales and inventory data in Excel to identify slow-moving SKUs and adjust reorder points, which reduced dead stock and improved cash flow.
- Studied seasonal demand patterns and vendor lead times to plan purchase cycles, helping balance stock availability against overstocking risk.
- Coordinated with suppliers on pricing and delivery schedules, working to keep procurement costs and inventory turnover in balance.

- Self-taught Python and SQL to convert manual sales-tracking spreadsheets into structured, repeatable reports, a hands-on introduction to data analysis that led directly to a formal transition into a Data Analyst role.

PROJECTS

SQIPE: Enterprise Analytics & ML Pricing Platform | Internal Project

Python | Scikit-learn | PostgreSQL | MySQL | dbt | FastAPI | Docker | Power BI

- Wrote ETL scripts to pull data from CRM webhooks and Tally transaction logs into one structured relational warehouse, consolidating pricing and sales data that used to live in scattered spreadsheets.
- Built pricing models that factor in non-linear component costs, switchgear configurations, and historical win-loss ratios, turning a manual, inconsistent pricing process into a standardized one.
- Deployed the model as an inference service with FastAPI across multiple Docker containers, with GitHub Actions handling CI/CD so updates ship without manual steps.
- Built Power BI dashboards backed by SQL to give leadership real-time visibility into pipeline risk and throughput.

Shoppers Intent MLOps | Personal Project

Python | FastAPI | MySQL | PostgreSQL | Scikit-learn | XGBoost | LightGBM | MLflow | Evidently AI | Docker | Render

- Trained and compared Random Forest, XGBoost, and LightGBM models to predict whether an online shopper would complete a purchase, tuning each with GridSearchCV. The best model reached **91.8%** accuracy and a **97.3%** ROC-AUC across **122,330** sessions.
- Corrected heavy class imbalance toward non-purchases using SMOTE along with feature scaling and selection, which raised recall to **94.3%** and meant far fewer actual buyers were missed.
- Used MLflow to track experiments and version models, then wrapped the final model in a FastAPI service containerized with Docker, with GitHub Actions handling automated testing and deployment on every push.
- Added drift monitoring with Evidently AI and a daily retraining job, and deployed the whole system live on Render instead of leaving it as a notebook.

OpinionMetric: AI Sentiment Transformer | Personal Project

Python | NLP | Scikit-learn | Streamlit | MySQL | Git | GitHub

- Built an NLP sentiment classifier on **50,033** labeled customer reviews, tagging each as positive, negative, or neutral with a confidence score, and reached **96%** accuracy with a **0.96** weighted F1-score.
- Wrote a text-cleaning pipeline that strips HTML, URLs, emojis, and punctuation before tokenizing, removing stopwords, and lemmatizing, then fed the cleaned text into a TF-IDF vectorizer and classifier.
- Logged every review, its predicted sentiment, and confidence score to a MySQL database, making it easy to audit predictions or retrain on real usage data later.
- Deployed the app on Streamlit Community Cloud with a clean interface and proper secrets management so database credentials stayed out of the codebase.

EDUCATION

B.E. in Computer Science & Engineering

2012 - 2016

Sipna College of Engineering, Amravati

[Prepared for competitive government exams (banking and State services)]

2016 - 2020

CERTIFICATIONS

- Diploma in Machine Learning using Python, Alison
- Data Science Essentials with Python, Cisco Networking Academy