

NIHAR GUDADHE

Data Scientist | AI ML Engineer

Nagpur, India | +91- 9168844536 | gudadhenihar@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Scientist & Analyst with 2 year 7 months of experience across the full machine learning lifecycle - EDA, feature engineering, model development, and deployment - for regression, classification, and time-series problems. Proficient in Python and SQL with a strong foundation in statistics. Productionized a Random Forest pricing model that cuts quotation turnaround from days to seconds and built automated CRM/ERP data pipelines to streamline the Quote-to-Cash process. Actively building skills in Generative AI (LLMs, prompt engineering, RAG) alongside core ML, MLOps, and Power BI analytics work.

TECHNICAL SKILLS

Programming & Databases: Python, SQL, PostgreSQL, MySQL, Excel

Machine Learning & Statistics: Regression, Classification, Clustering, Time Series Forecasting, Ensemble Methods (Random Forest, XGBoost, LightGBM), PCA, Anomaly Detection, Feature Engineering & Selection, Cross-Validation, Hyperparameter Tuning, SMOTE, Hypothesis Testing, A/B Testing

Deep Learning & NLP: ANN, CNN, RNN, LSTM, GRU, Encoder-Decoder, Attention Mechanism, Tokenization, Word Embeddings, Named Entity Recognition, Sentiment Analysis

Generative AI: Large Language Models (LLMs), Prompt Engineering, RAG, Hugging Face Transformers, LangChain

Computer Vision: Image Classification, Object Detection (YOLO), Image Segmentation, OpenCV

Libraries & Frameworks: Pandas, NumPy, Scikit-learn, TensorFlow, PyTorch, Matplotlib, Seaborn, NLTK, spaCy, PySpark,

Databricks Data Visualization: Power BI, DAX, Power BI Service, Tableau, Excel

Engineering & Deployment: FastAPI, Docker, Git, GitHub, CI/CD, AWS, Render, Jupyter Notebook, VS Code

Core Concepts: Statistical Inference, Confidence Intervals, EDA, Data Cleaning, Data Modeling, Data Warehousing, Data Labeling.

PROFESSIONAL EXPERIENCE

Data Scientist & Analyst - Influx Control Panel, Nagpur

Feb 2024 - Present

- **Enterprise Data Integration & Warehousing:** Architected an automated ELT pipeline that pulls raw CRM leads and billing records into a MySQL staging lake, transforms the data through dbt, and lands clean, analysis-ready tables in a PostgreSQL warehouse - replacing what used to be manual, error-prone data pulls.
- **Predictive Pricing Engine:** Built and productionized a Random Forest regression model in Scikit-learn to automate cost estimation for complex industrial items, deploying it behind a real-time FastAPI endpoint in Docker - cutting quotation turnaround from days to seconds.
- **Operational Efficiency:** Developed a production line efficiency dashboard tracking cycle times and machine idle time, which surfaced hidden bottlenecks on the shop floor and drove a **12%** increase in throughput.
- **Sales Performance Analytics:** Designed a **360-degree** customer dashboard in Power BI covering account value, engagement, and conversion stage, giving the sales team a clear way to prioritize high-value segments - contributing to a **10%** lift in conversion rates.
- **Quality Assurance:** Applied anomaly detection to build an automated defect - analysis report linking raw material batches to support ticket trends, helping engineering catch recurring issues early and cut product return rates by **14%**.

PROJECTS

SQIPE - Enterprise Analytics & ML Pricing Platform | Internal Project

Python | Scikit-learn | PostgreSQL | MySQL | dbt | FastAPI | Docker | Power BI

- Wrote automated ETL scripts to pull data from CRM webhooks and Tally transaction logs and sync it into a properly structured relational warehouse, so pricing and sales data lived in one consistent schema instead of scattered sources.
- Built feature-engineered pricing models that factor in non-linear component costs, switchgear configurations, and historical win-loss ratios, replacing inconsistent manual pricing with a standardized, data-driven approach.

- Deployed the model as an inference service using FastAPI across multi-container Docker environments, with GitHub Actions handling CI/CD so updates could ship without manual deployment steps.
- Built SQL-backed Power BI dashboards to give leadership real-time visibility into industrial pipeline risk and throughput, turning raw transaction data into something the team could actually act on.

Shoppers Intent MLOps | Personal Project

Python | FastAPI | MySQL | PostgreSQL | Scikit-learn | XGBoost | LightGBM | MLflow | Evidently AI | Docker | Render

- Trained and compared Random Forest, XGBoost, and LightGBM models to predict whether an online shopper would actually buy something, using GridSearchCV to tune each one - the best model hit **91.8%** accuracy and **97.3%** ROC-AUC on over 122,330 sessions.
- Noticed the data was heavily skewed toward non-purchases, so I applied SMOTE along with feature scaling and selection to fix the imbalance - this pushed recall up to **94.3%**, meaning far fewer actual buyers were getting missed by the model.
- Set up MLflow to track every experiment and version each model, then wrapped the final model in a FastAPI service and containerized it with Docker, with GitHub Actions handling automated testing and deployment on every code push.
- Added drift monitoring with Evidently AI and a daily auto-retraining job so the model wouldn't go stale over time - deployed the whole thing live on Render instead of leaving it as a notebook.

Opinion Metrix: AI Sentiment Transformer | Personal Project

Python | NLP | Scikit-learn | Streamlit | MySQL | Git | GitHub

- Built an NLP sentiment classifier trained on over 50,000+ labeled customer reviews to tag each one as Positive, Negative, or Neutral along with a confidence score, landing at **96%** accuracy and a **0.96** weighted F1-score.
- Wrote a full text-cleaning pipeline that strips out HTML, URLs, emojis, and punctuation before tokenizing, removing stopwords, and lemmatizing - then fed the cleaned text into a TF-IDF vectorizer and ML classifier for prediction.
- Connected the app to a MySQL database so every review, its predicted sentiment, and the model's confidence score gets logged automatically, making it easy to audit predictions later or retrain on real usage data.
- Deployed the whole app live on Streamlit Community Cloud with a clean, interactive UI and proper secrets management so database credentials never got exposed in the codebase.

EDUCATION

B.E. in Computer Science Engineering

Sipna College of Engineering, Amravati

CERTIFICATIONS

- Diploma in Machine Learning using Python - Alison
- Data Science Essentials with Python - Cisco Networking Academy