

NIHAR GUDADHE

Data Scientist | AI ML Engineer

Nagpur, India | +91- 9168844536 | gudadhenihar@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Scientist & Analyst with 2 years 7 months of experience across the full machine learning lifecycle - EDA, feature engineering, model development, and deployment - for regression, classification, and time-series problems. Proficient in Python and SQL with a strong foundation in statistics. Productionized a Random Forest pricing model that cuts quotation turnaround from days to seconds and built automated CRM/ERP data pipelines to streamline the Quote-to-Cash process. Actively building skills in Generative AI (LLMs, prompt engineering, RAG) alongside core ML, MLOps, and Power BI analytics work.

TECHNICAL SKILLS

Programming & Databases: Python, SQL, PostgreSQL, MySQL, Excel

Machine Learning & Statistics: Regression, Classification, Clustering, Time Series Forecasting, Ensemble Methods (Random Forest, XGBoost, LightGBM), PCA, Anomaly Detection, Feature Engineering & Selection, Cross-Validation, Hyperparameter Tuning, SMOTE, Hypothesis Testing, A/B Testing

Deep Learning & NLP: ANN, CNN, RNN, LSTM, GRU, Encoder-Decoder, Attention Mechanism, Tokenization, Word Embeddings, Named Entity Recognition, Sentiment Analysis

Generative AI: Large Language Models (LLMs), Prompt Engineering, RAG, Hugging Face Transformers, LangChain

Computer Vision: Image Classification, Object Detection (YOLO), Image Segmentation, OpenCV

Libraries & Frameworks: Pandas, NumPy, Scikit-learn, TensorFlow, PyTorch, Matplotlib, Seaborn, NLTK, spaCy, PySpark, Databricks

Data Visualization: Power BI, DAX, Power BI Service, Tableau, Excel

Engineering & Deployment: FastAPI, Docker, Git, GitHub, CI/CD, AWS, Render, Jupyter Notebook, VS Code

Core Concepts: Statistical Inference, Confidence Intervals, EDA, Data Cleaning, Data Modeling, Data Warehousing, Data Labeling

PROFESSIONAL EXPERIENCE

Data Scientist & Analyst - Influx Control Panel, Nagpur

Feb 2024 - Present

- Enterprise Data Integration & Warehousing:** Architected an automated ELT pipeline pulling raw CRM leads and billing records into a MySQL staging lake, transforming the data through dbt, and loading clean, analysis-ready tables into a PostgreSQL warehouse. This replaced manual, error-prone data pulls the team previously relied on.
- Predictive Pricing Engine:** Built and productionized a Random Forest regression model in Scikit-learn to automate cost estimation for complex industrial items. Deployed behind a real-time FastAPI endpoint running in Docker, it cut quotation turnaround from days to seconds.
- Operational Efficiency:** Developed a production line efficiency dashboard to track cycle times and machine idle time. The bottlenecks it surfaced on the shop floor led directly to a **12%** increase in throughput.
- Sales Performance Analytics:** Designed a **360-degree** customer dashboard in Power BI covering account value, engagement, and conversion stage. Sales used it to prioritize high-value segments, and conversion rates improved by **10%** as a result.
- Quality Assurance:** Applied anomaly detection to build an automated defect-analysis report correlating raw material batches with support ticket trends. Engineering used the findings to catch recurring issues earlier, cutting product return rates by **14%**.

PROJECTS

SQIPE - Enterprise Analytics & ML Pricing Platform | Internal Project

Python | Scikit-learn | PostgreSQL | MySQL | dbt | FastAPI | Docker | Power BI

- Wrote automated ETL scripts pulling data from CRM webhooks and Tally transaction logs into a single, properly structured relational warehouse, consolidating pricing and sales data that had previously lived in scattered sources.
- Built feature-engineered pricing models factoring in non-linear component costs, switchgear configurations, and historical win-loss ratios, standardizing a pricing process that used to be manual and inconsistent.

- Deployed the model as an inference service using FastAPI across multi-container Docker environments; GitHub Actions handled CI/CD so updates shipped without manual steps.
- Built SQL-backed Power BI dashboards giving leadership real-time visibility into industrial pipeline risk and throughput.

Shoppers Intent MLOps | Personal Project

Python | FastAPI | MySQL | PostgreSQL | Scikit-learn | XGBoost | LightGBM | MLflow | Evidently AI | Docker | Render

- Trained and compared Random Forest, XGBoost, and LightGBM models to predict whether an online shopper would actually complete a purchase, tuning each with GridSearchCV. The best model reached **91.8%** accuracy and **97.3%** ROC-AUC across **122,330** sessions.
- The data was heavily skewed toward non-purchases, so SMOTE was applied along with feature scaling and selection to correct the imbalance - recall rose to **94.3%**, meaning far fewer actual buyers were missed by the model.
- Set up MLflow to track every experiment and version each model, then wrapped the final model in a FastAPI service containerized with Docker. GitHub Actions handled automated testing and deployment on every push.
- Added drift monitoring with Evidently AI and a daily auto-retraining job to keep the model from going stale and deployed the full system live on Render rather than leaving it as a notebook.

OpinionMetric: AI Sentiment Transformer | Personal Project

Python | NLP | Scikit-learn | Streamlit | MySQL | Git | GitHub

- Built an NLP sentiment classifier trained on **50,033** labeled customer reviews to tag each one as **Positive**, **Negative**, or **Neutral** along with a confidence score, reaching **96%** accuracy and a **0.96** weighted F1-score.
- Wrote a full text-cleaning pipeline that strips out HTML, URLs, emojis, and punctuation before tokenizing, removing stopwords, and lemmatizing, then fed the cleaned text into a TF-IDF vectorizer and ML classifier for prediction.
- Connected the app to a MySQL database so every review, its predicted sentiment, and the model's confidence score gets logged automatically, making it easy to audit predictions later or retrain on real usage data.
- Deployed the app live on Streamlit Community Cloud with a clean, interactive UI and proper secrets management so database credentials never got exposed in the codebase.

EDUCATION

B.E. in Computer Science Engineering

Sipna College of Engineering, Amravati

CERTIFICATIONS

- **Diploma in Machine Learning using Python** - Alison
- **Data Science Essentials with Python** - Cisco Networking Academy